

Data Analytics for Business

Nonlinear Transformation Models

Sridhar Narasimhan, Ph.D

Professor

Scheller College of Business

Introduction to Nonlinear Models




Lessons in this Module

- A. Introduction to Nonlinear Models
- B. Linear-Log Model
- C. Log-Linear Model
- D. Log-Log Model
- E. Polynomial Model

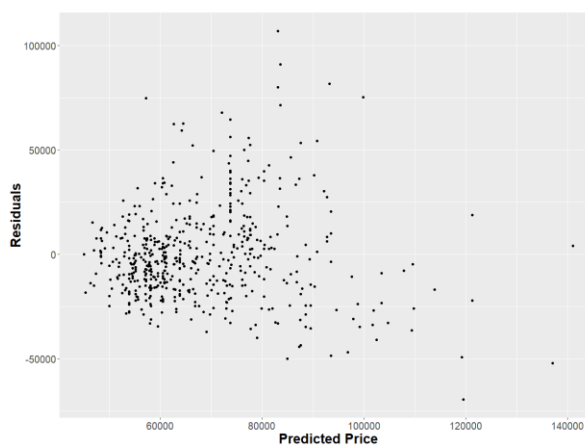


An Example of a Nonlinear Relationship



Georgia
Tech

Another Example (From the Housing Dataset in the Ecdat Package in R)



Georgia
Tech


Model A: Linear-Linear Model

	Estimate	S.E.	t Value	Pr> t
Intercept	34140	2,491	13.7	<.0001
lotsize	6.599	0.4458	14.8	<.0001

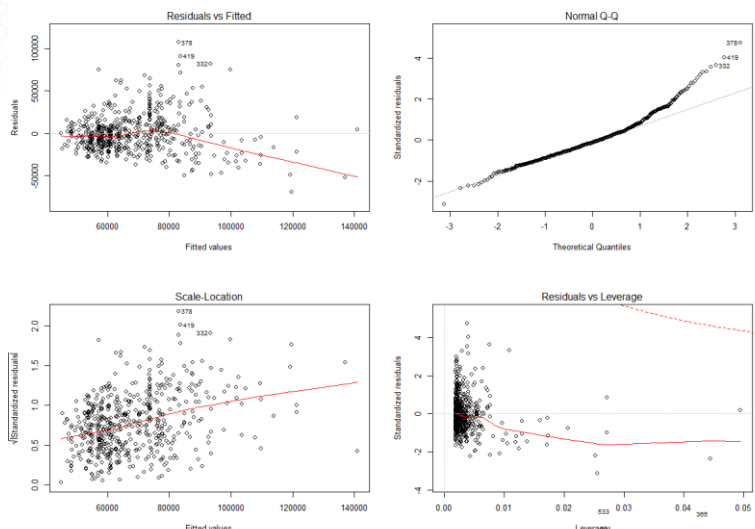
R-Squared	Adjusted R-Squared
0.2871	0.2858

- Model A: $price = b_0 + b_1 * lotsize$
- As X (lotsize) increases by 1 unit, Y (price) changes by b_1 (6.599) units, holding all other factors constant

Model A: Scatter Plot



Model A: Diagnostics Plot



Georgia
Tech

Model A: $price = b_0 + b_1 * lotsize$ Assumptions

- We need to check if the assumptions for a linear regression model hold
- The residuals vs. fitted values plot indicates that Model A has heteroscedasticity (non-constant variance); the QQ-plot suggest that there is non-linearity
- Hence, we need to start exploring non-linear models
- We focus on natural log transformations of the variables because they are easier to interpret. Note that in R, `log()` computes the natural log

Georgia
Tech

The Various Log Transformations

	Y	log(Y)
X	Model A: Level-level model $Y = b_0 + b_1 * X$	
log(x)		

The Various Log Transformations

	Y	log(Y)
X	Model A: Level-level model $Y = b_0 + b_1 * X$	
log(x)	Model B: Linear-Log Model $Y = b_0 + b_1 * \log(X)$	

The Various Log Transformations

	Y	log(Y)
X	<i>Model A: Level-level model</i> $Y = b_0 + b_1 * X$	Model C: Log-linear Model $\log(Y) = b_0 + b_1 * X$
log(x)	<i>Model B: Linear-Log Model</i> $Y = b_0 + b_1 * \log(X)$	

The Various Log Transformations

	Y	log(Y)
X	<i>Model A: Level-level model</i> $Y = b_0 + b_1 * X$	<i>Model C: Log-linear Model</i> $\log(Y) = b_0 + b_1 * X$
log(x)	<i>Model B: Linear-Log Model</i> $Y = b_0 + b_1 * \log(X)$	Model D: Log-Log $\log(Y) = b_0 + b_1 * \log(X)$

Note:

- The log() function in R computes the natural logarithm
- If a variable x has some values = 0, then use log(x+1) transformation

Data Analytics for Business

Nonlinear Transformation Models

Sridhar Narasimhan, Ph.D

Professor

Scheller College of Business

Linear-Log Model

Georgia
Tech

Model B: Linear-Log Model Independent Variable Transformed

	Estimate	S.E.	t Value	Pr> t
Intercept	-250,728	20,184	-12.42	<.0001
Ln_lotsize	37660	2381	15.81	<.0001

R-Squared	Adjusted R-Squared
0.315	0.3137

- Model B: $price = b_0 + b_1 \cdot \log(lotsize)$
- Create a new variable Ln_lotsize which is the natural log of lotsize
- Run Model B using the Housing dataset
- How would you interpret the coefficient of Ln_lotsize?

Georgia
Tech

Model B: $price = b_0 + b_1 * \log(lotsize)$

What does the coefficient $b_1 = 37660$ imply?

- A. When lotsize increases by 1 sq foot, price increases (on average) by \$37,660.
- B. When lotsize decreases by 1 sq foot, price increases (on average) by \$37,660.
- C. When lotsize decreases by $\log(1)$, price increases (on average) by \$376.60.
- D. When lotsize increases by 1%, price increases (on average) by \$376.60.

What is the correct answer?

D. A one percent increase in the independent variable increases (or decreases) the dependent variable by (coefficient/100) units.




Interpreting a Linear-Log Model

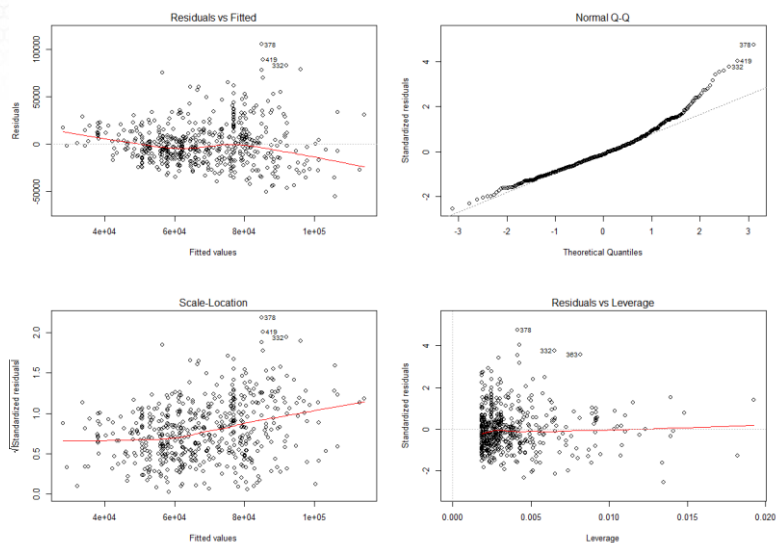
- Linear-log model (and most other models) needs to be interpreted carefully
- It does not make much practical sense to increase the “ $\log(\text{lotprice})$ ” by one unit
- But increasing X by 1 percent is almost equivalent to increasing (natural) $\log(X)$ by 0.01 units
- Hence, a 1 percent increase in X increases (natural) $\log(X)$ by .01 and, therefore, changes the Y variable by $.01 * b_1$



Model B: Scatter Plot



Model B: Diagnostics Plots



Data Analytics for Business

Nonlinear Transformation Models

Sridhar Narasimhan, Ph.D

Professor

Scheller College of Business

Log-Linear Model

Georgia
Tech

Model C: Log-Linear Model Dependent Variable Transformed

	Estimate	S.E.	t Value	Pr> t
Intercept	10.58	.03451	306.51	<.0001
lotsize	.00009315	.000006177	15.08	<.0001

R-Squared	Adjusted R-Squared
0.2947	0.2935

- Model C: $\log(\text{price}) = b_0 + b_1 * \text{lotsize}$
- Create a new variable Ln_price which is the natural log of price
- How would you interpret the coefficient of lotsize?

Georgia
Tech

Model C: $\log(\text{price}) = b_0 + b_1 * \text{lotsize}$

What does the coefficient $b_1 = .00009315$ imply in a log-linear model?

- A. When lotsize increases by .01 sq. ft., price increases by .009315% (on average)
- B. When lotsize increases by 1 sq. ft., price increases by .009315% (on average)
- C. When lotsize increases by 1 sq. ft., price decreases by .009315% (on average)
- D. When lotsize decreases by 1 sq. ft., price increases by .009315% (on average)

What is the correct answer?

B. The dependent variable changes by $100 * (\text{coefficient})$ percent for a one unit increase in the independent variable while all other variables in the model are held constant




Interpreting the Log-linear Model

$$\log(\text{price}) = b_0 + b_1 * \text{lotsize}$$

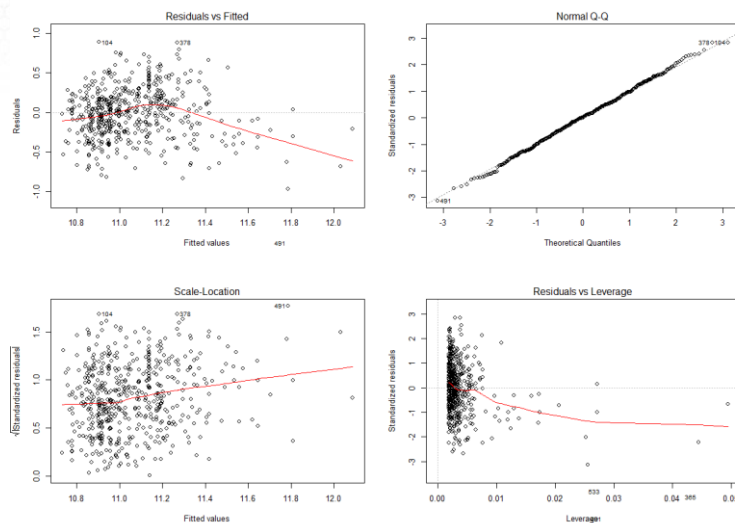
- Increasing x by one unit will increase (natural) $\log(y)$ by b_1 units
- With $x = \text{lotsize}$, the model $\log(\text{price}) = b_0 + b_1 * x$ is the same as $y = e^{(b_0 + b_1 x)}$. Hence, $dy/dx = b_1 y$, or $dy/y = b_1 * dx$
- Multiplying both sides by 100, we get $100 * dy/y = 100 * b_1 * dx$
- Note that $(100 * dy/y)$ is the percentage change in Y
- If $dx = 1$, then this one unit change in x leads to a $100 * b_1$ percentage change in Y
- Note: this interpretation works when $b_0 + b_1 * x$ is very small
The accurate percentage change in $Y = (e^{b_1} - 1) * 100$ for a one unit change in X



Model C: Scatter Plot



Model C: Diagnostic Plots



Data Analytics for Business

Nonlinear Transformation Models

Sridhar Narasimhan, Ph.D

Professor

Scheller College of Business

Log-Log Model

Georgia
Tech

Model D: Log-Log Model

	Estimate	S.E.	t Value	Pr> t
Intercept	6.46853	0.27674	23.37	<.001
lotsize	0.54218	0.03265	16.61	<.001

R-Squared	Adjusted R-Squared
0.3364	0.3352

- Model D: $\log(\text{price}) = b_0 + b_1 * \log(\text{lotsize})$
- The dependent and the independent variables are log transformed
- How do you interpret the coefficient of lotsize?

Georgia
Tech

Model D: Log-Log Model

$$\log(\text{price}) = b_0 + b_1 * \log(\text{lotsize})$$

$b_1 = 0.54218$ implies that

- A. When lotsize increases by 1%, price increases (on average) by 0.54218%
- B. When lotsize decreases by 1%, price increases (on average) by 0.54218%
- C. When lotsize increases by 1 sq. ft., price increases (on average) by 0.54218%
- D. When lotsize decreases by 1 sq. ft., price increases (on average) by 0.54218%

What is the correct answer?

Answer -- A. *The dependent variable changes by b_1 % percent for a one percent increase in the independent variable while all other variables in the model are held constant*




Interpreting the Log-Log Model

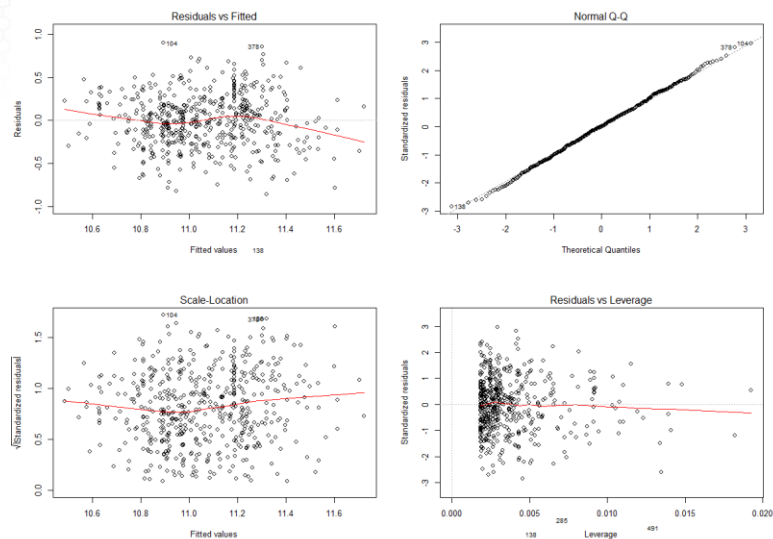
- Increasing (natural) $\log(X)$ by 0.01 leads to increasing (natural) $\log(Y)$ by $b_1 * 0.01$ units
- Increasing (natural) $\log(X)$ by 0.01 is almost equivalent to increasing X by 1 percent, which implies changing Y by b_1 percent
- In a regression setting, we'd interpret elasticity as the percent change in Y (the dependent variable), when X (the independent variable) increases by one percent
- Hence b_1 captures elasticity



Model D: Scatter Plot



Model D: Diagnostics Plot



Comparing the Models

	R-Squared	Adjusted R-Squared
Model A: Level-Level	0.2871	0.2858
Model B: Linear-Log	0.3150	0.3137
Model C: Log-Linear	0.2947	0.2935
Model D: Log-Log	0.3364	0.3352

Reasons for (log) Transforming Data

- To achieve a (more) linear relationship
- To make a distribution more normal
- To make the variance more constant
- To get a better fit in the model – i.e., increase R-Squared

Log Transformations Cheat Sheet

	Y	log(Y)
X	Model A: Level-Level model $Y = b_0 + b_1 * X$ As X increases by 1 unit, Y changes by b_1 units, holding all other factors constant.	Model C: Log-Linear Model $\log(Y) = b_0 + b_1 * X$ As X increases by 1 unit, Y increases by $(b_1 * 100)\%$, holding all other factors constant.
log(x)	Model B: Linear-Log Model $Y = b_0 + b_1 * \log(X)$ As X increases by 1%, Y increases by $(b_1/100)$ units, holding all other factors constant.	Model D: Log-Log $\log(Y) = b_0 + b_1 * \log(X)$ As X increases by 1%, Y changes by $b_1\%$, holding all other factors constant.

Data Analytics for Business

Nonlinear Transformation Models

Sridhar Narasimhan, Ph.D

Professor

Scheller College of Business

Polynomial (Quadratic) Model

Model E: Polynomial (Quadratic) Model

	Estimate	S.E.	t Value	Pr> t
Intercept	11,340	4,892	2.317	.0209*
lotsize	14.81	1.589	9.317	<.0001
lot_square	-6.238e-04	1.162e-04	-5.370	<.0001

R-Squared	Adjusted R-Squared
0.323	0.3205

- Model E: $price = b_0 + b_1 * lotsize + b_2 * lotsize^2$
- Create a new variable $lot_square = lotsize^2$
- Fit the model, $price = b_0 + b_1 * lotsize + b_2 * lotsize^2$



Model E: $price = b_0 + b_1 * lotsize + b_2 * lotsize^2$

The coefficient $b_1 = 14.81$ implies that

- When lotsize increases by 1 sq. ft., price will increase by \$14.81
- When lotsize decreases by 1 sq. ft., price will increase by \$14.81
- When lotsize increases by 1%, price will increase by \$14.81
- None of the above

What is the correct answer?

D: None of the above



Model E: Interpreting a Polynomial (Quadratic) Model

$$\text{price} = b_0 + b_1 * \text{lotsize} + b_2 * \text{lotsize}^2$$

- Coefficients b_1 and b_2 cannot be interpreted individually because when *lotsize* is increased by 1 unit, it is not possible (or meaningful) to hold *lotsize*² constant
- A quadratic model does not allow for an isolated interpretation of coefficients since $\frac{d(\text{price})}{d(\text{lotsize})} = b_1 + 2b_2 * \text{lotsize}$
- This means that the slope (impact of 1 unit increase in x) is not a constant. It changes at every point of the quadratic curve (if you plot y vs. x)



Recap of this Module

- Introduction to Nonlinear Models
- Linear-Log Model
- Log-Linear Model
- Log-Log Model
- Polynomial Model

